

REVIEW	CODE REVIEW	HISTORY
<h2>Meets Specifications</h2> <p>Dear Udacian, the project is very well implemented and meets the specifications! Congratulations on successfully completing the project. As a next step, please go through the resource: <a href="#">human-like robot hand trained to manipulate physical objects with unprecedented dexterity</a>.</p> <p>All the best for the next projects! 😊</p>		
<h2>Training Code</h2> <div> <div>✓</div> <div> <p>The repository (or zip file) includes functional, well-documented, and organized code for training the agent.</p> <p>The repository contains jupyter notebook, code files, readme, and project report. The code is functional. The implemented Deep Q Network and Double Deep Q Network look good!</p> <p>You should definitely check the following resource to progress with the learning further:</p> <ul style="list-style-type: none"> <li>• <a href="#">Speeding up DQN on PyTorch: how to solve Pong in 30 minutes</a></li> <li>• <a href="#">Advanced DQNs: Playing Pac-man with Deep Reinforcement Learning by mapping pixel images to Q values</a></li> </ul> </div> </div>		
<div> <div>✓</div> <div> <p>The code is written in PyTorch and Python 3.</p> <p>The code is written in pytorch and python3.</p> <ul style="list-style-type: none"> <li>• A good read: <a href="#">PyTorch vs TensorFlow—spotting the difference</a></li> </ul> </div> </div>		
<div> <div>✓</div> <div> <p>The submission includes the saved model weights of the successful agent.</p> <p>Thanks for including the saved model weights of the successful agent.</p> </div> </div>		
<h2>README</h2> <div> <div>✓</div> <div> <p>The GitHub (or zip file) submission includes a <code>README.md</code> file in the root of the repository.</p> </div> </div>		
<div> <div>✓</div> <div> <p>The README describes the the project environment details (i.e., the state and action spaces, and when the environment is considered solved).</p> <p>Awesome work providing the project environment details in the README. State space, action space, reward function and when the environment is considered solved is specified very informatively.</p> </div> </div>		
<div> <div>✓</div> <div> <p>The README has instructions for installing dependencies or downloading needed files.</p> <p>Proper instructions have been specified in the README to install the dependencies and download the necessary files.</p> </div> </div>		
<div> <div>✓</div> <div> <p>The README describes how to run the code in the repository, to train the agent. For additional resources on creating READMEs or using Markdown, see <a href="#">here</a> and <a href="#">here</a>.</p> </div> </div>		

## Report

✓	<p>The submission includes a file in the root of the GitHub repository or zip file (one of <code>Report.md</code>, <code>Report.ipynb</code>, or <code>Report.pdf</code>) that provides a description of the implementation.</p>
	<p>Report.md has been included in the root of the github repository.</p>
✓	<p>The report clearly describes the learning algorithm, along with the chosen hyperparameters. It also describes the model architectures for any neural networks.</p>
	<p>The report is rather informative providing an insight on every aspect of the project which includes Implementation, model architectures, hyperparameters, rewards, future works.</p> <ul style="list-style-type: none"><li>• Good implementation of the agent for Deep Q Network and Double Deep Q Network.</li><li>• Correct decoupling of the parameters being updated from the ones that are using a target network to produce target values.</li><li>• Perfect implementation of The Epsilon-greedy action selection to encourage exploratory behavior in the agent.</li><li>• Good use of tau parameter to perform soft-update. It helps to prevent variance into the process due to individual batches.</li><li>• Good use of the replay memory to store and recall experience tuples.</li><li>• The implementation is easy to debug and easily extensible.</li></ul>
✓	<p>A plot of rewards per episode is included to illustrate that the agent is able to receive an average reward (over 100 episodes) of at least +13. The submission reports the number of episodes needed to solve the environment.</p>
	<p><b>Awesome</b></p> <ul style="list-style-type: none"><li>• The agent seems to perform very well!</li><li>• The agent is able to achieve a reward of 13+ over last 100 episodes in just<ul style="list-style-type: none"><li>◦ 464 episodes for DQN!</li><li>◦ 356 episodes for DDQN!</li></ul></li><li>• The submission discusses the rewards plot obtained clearly.</li></ul>
✓	<p>The submission has concrete future ideas for improving the agent's performance.</p>
	<p>Great job providing the ideas to experiment more in future with the project!</p> <p>I would like to point you to the following resources:</p> <ul style="list-style-type: none"><li>• <a href="#">Rainbow: Combining Improvements in Deep Reinforcement Learning</a></li><li>• <a href="#">Conquering OpenAI Retro Contest 2: Demystifying Rainbow Baseline</a></li></ul> <p>As pointed in the report, you should try implementing Prioritized Experience Replay also. It helps to improve the performance and significantly reduces the training time. A fast implementation of Prioritized Experience Replay is possible using a special data structure called a Sum Tree. I found a <a href="#">good implementation here</a>.</p> <p>And, you should definitely try applying these algorithms by taking raw screen pixels as input also. In case you get stuck, definitely check <a href="#">this github repository</a>.</p>